Methods Addendum v1.1

To *Watching AI Become an Educator: MIBO-Education and the Longitudinal Observation of Pedagogical Drift*

Kento Sasano | ORCID 0009-0009-3853-8029
Okayama University; Keio Research Institute at SFC, Keio University

Adopted 12 November 2026. Protocol DOI: **10.5281/zenodo.24157241**
English | CC BY 4.0 | Existing SSRN paper: 10.2139/ssrn.7330038

1. Prospective adoption and scope

This addendum prospectively governs Education W01 v1.1. It supersedes the old five-series target population, human/fixed-panel rater assignment, elapsed 1 September baseline and technical-retry provisions for this future wave only. The unchanged v1.0 manuscript follows this addendum as a dated historical source. Original protocol DOI 10.5281/zenodo.22047001, old ZIPs, frozen scientific materials, W0 data, judgments and machine flags remain preserved. W0 is nonofficial engineering evidence, not the formal baseline. No W01 results exist at adoption.

The owner explicitly approved the exact reviewed content snapshot. Its preparation-stage DRAFT labels are preserved in the unchanged ZIP; a separate dated finalization declaration adopts those scientific-content bytes as formal v1.1. Actual Zenodo publication must be independently verified before model execution. This addendum updates existing SSRN paper 7330038, not a duplicate new study.

2. Design, model identities and reciprocal assignment

Observe 70 frozen Japanese EBB-JA v1.0 items with 10 independent single-turn replications in each of three first-party CLOSED API series: 2,100 planned source slots. Each available source response is scored by the two nonself provider families. Six exact requested model IDs are frozen; fresh metadata checks establish access/identity evidence, not unchanged underlying model weights. No silent substitution or upgrade is permitted.

Source series	Requested source model	Nonself judge IDs
M01 OpenAI	gpt-5.6-sol	J01 + J02
M02 Anthropic	claude-opus-5	J02 + J03
M03 Gemini	gemini-3.6-flash	J01 + J03

Judges: J01 Anthropic **claude-opus-5-5**; J02 Gemini **gemini-3.8-flash**; J03 OpenAI **gpt-6.1-sol**. Source generation cap: 8,192. Judge caps: J01 8,192; J02 16,384; J03 8,192. The W0-tested v3.1 scoring prompt, item anchors, output contract and aggregation remain fixed.

Source requests retain exact frozen item text and the W0-tested payload structure: no added source system/developer prompt, conversational history, tools, web, files or memory; native-default sampling. Retain exact wire hashes and actual request/response times. The logical observation site is MIBO-SITE-SASANO-WINDOWS-01, timezone Asia/Tokyo; geographic region has not been independently verified.

3. Windows, caps and one-attempt policy

W01 baseline: **20 November 2026, 09:00 JST (00:00 UTC)**. Fresh metadata and common calibration run in the preceding 24 hours, starting no earlier than 19 November 09:00 JST. Collection targets 12 hours with a hard 24-hour limit (21 November 09:00 JST). Judging ends within 72 hours after actual collection completion. No automatic reset or extension of dates/windows is authorized.

Maximum new reservations: **6,411** = 6 metadata checks + 105 calibration ratings (35 examples x 3 judges) + 2,100 source slots + 4,200 nonself judging slots. Preserve 1,468 prior reservations; cumulative ceiling **7,879**. Caps are not guaranteed successful sample sizes or a USD cost guarantee. The earlier USD 5 condition was withdrawn; it is not reinstated here.

The prospective attempt limit changes from at most five technical attempts to **one attempt per slot, zero automatic resends**. A fully recorded READ_TIMEOUT consumes a reservation and remains technical missingness; only unsent slots may continue. Missing source cells are not refilled, and their two judging requests are not sent. Unknown/unrecorded outcomes, authentication, HTTP, identity or CLOSED-control failures trigger HOLD. Scientific refusal, blocking, truncation, poor answers and invalid model outputs are data, not reasons to repair or retry.

Dispatch remains serial with at least 2 seconds between requests and a 600-second communication timeout. Existing approved credentials and the office-managed paid Gemini project are retained. During actual AC-connected execution only, SystemRequired/ExecutionRequired sleep inhibition is held and is cleared/closed on normal exit, error or AC loss. No permanent power-setting change or explicit-sleep override is authorized. Core and all other waves are excluded.

4. Common calibration and delegated technical review

Repeat the approved common 35-example calibration bank for all three judges before observation. Preserve all individual judgments, disagreements and missingness. Apply the fixed separate primary-score and full-record numerical qualification guards and an explicit delegated AGENT definition/adversarial review; numerical passage alone does not authorize collection. The v3.1 developmental four-item 0-2 definitions and refusal-absent R0 rule remain unchanged. Do not repair Gemini invalid output or silently promote W0 failed flags.

Primary scoring and all 45 auxiliary diagnostics remain separate. Agreement between AI judges, including common calibration, is not human gold-standard validity, independence, all-45-field semantic validation or cross-source comparability. A rejected calibration is preserved and stops W01; model, prompt, anchors or criteria cannot be changed to manufacture passage.

5. Analysis and interpretation

Retain original Codebook section 3 arithmetic: each within-source domain index is the mean of valid item scores in that domain; a 0-100 display is 25 times that mean. There are 100 planned slots per domain/source. Retain all replications and report valid denominators, unresolved disagreements, incomplete/missing counts, distributions, medians and descriptive dispersion. Do not apply W0's sample-150/50 readiness denominator to the formal 2,100-slot wave.

No overall best-AI ranking, pooled changing-rater kappa, independence-based confidence interval, causal claim or W0-as-official-baseline comparison is introduced. W01 establishes a baseline, not longitudinal drift. Source-pair diagnostics and auxiliary fields remain separate. Identical reported model IDs do not prove stable model weights or strengthen a comparability classification.

6. Availability, authorship and unchanged disclosures

Protocol materials: https://doi.org/10.5281/zenodo.24157241. Previous version: 10.5281/zenodo.22047001. The new DOI was actually reserved before finalization; publication and file-byte verification are separate operational receipts. Pending SSRN review is not described as public availability. Protocol materials remain CC BY 4.0; source code retains its existing license. This machine-only protocol does not recruit human participants; no institutional ethics approval is invented.

AI-assistance disclosure. Generative-AI systems assisted with literature discovery, research organization, language editing, document production, structural critique, and consistency checking. The author selected the research design, verified cited sources and project records, made all substantive decisions, and accepts responsibility for the manuscript. Generative-AI systems are not authors.

Competing interests. Kento Sasano is the chief executive officer and an equity holder of Scrumy Inc., a company operating in AI governance, AI security, and AI marketing. He also serves as an AI strategy advisor to Booost Inc. These relationships are disclosed because their activities may be perceived as relevant to AI evaluation and governance. The author retains responsibility for the research design, interpretation, and manuscript. No other competing interests are declared.

Funding. This research received no external financial or in-kind support and was self-funded by the author. Scrumy Inc. and Booost Inc. provided no financial or in-kind support for this research.

Additional methods disclosure. AI systems act as the reciprocal scorers described above. An AI agent performs only the expressly delegated limited calibration/results technical review. The author retains substantive responsibility; AI systems are not independent human raters or authors.

7. Content identity and historical-manuscript boundary

Approved preparation release SHA-256:

`64c257ffaf13e447dca643275fa46e8dcdf7beef9ab100a4ed1fb7fc26438f36`

Unchanged adopted ZIP SHA-256:

`4854f94fb63fe7d63a7d318a45372c9daeae718c696b724239c76f2c17398da3`

The following v1.0 pages are retained unchanged in content. Where they conflict with this prospective addendum, this addendum governs W01 v1.1 only. Historical dates, original design descriptions and preliminary hypotheses in the retained manuscript are not new observations or results.